

# Cervix Type Classification for Cancer Treatment with convolutional neural networks

*By: TheDoctors*

Iuliia Klykova  
Emmanuel Ojuba  
Victoria Gianello

# Problem description

Cervix type

Different transformation zone locations =  
Different Cervix type

Source: The Cervix, Singer et al, 2006

**Type 1**


- Completely ectocervical
- Fully visible
- Small or large

**Type 2**

- Has endocervical component
- Fully visible
- May have ectocervical component which may be small or large

**Type 3**

- Has endocervical component
- Is not fully visible
- May have ectocervical component which may be small or large



Cervical cancer is easy to prevent in its precancerous stage. However, many times doctors are not able to give patients the appropriate treatment for their cervix anatomy. Therefore, we created a Neural Network that given an image of a cervix, classifies cervix types accurately. This will allow doctors give the right treatment more reliably.

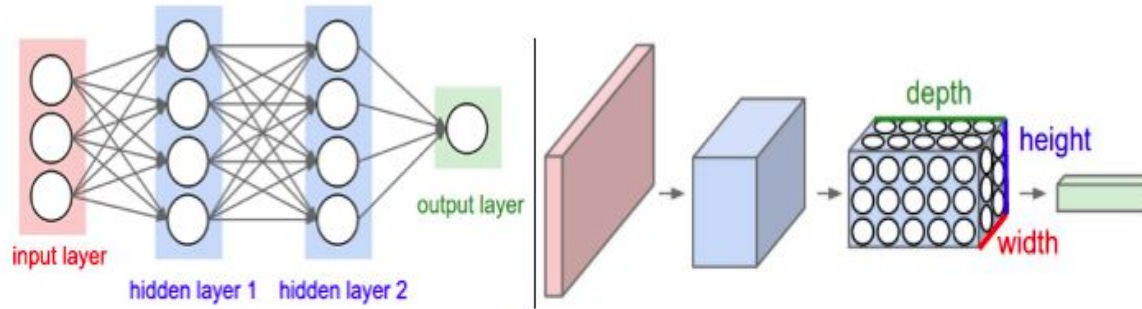
Initial Training DataSet ~ 1500 images



Figure 2. Example images of each cervix type. Image (a) is Type 1, image (b) is Type 2, and Image (c) is Type 3

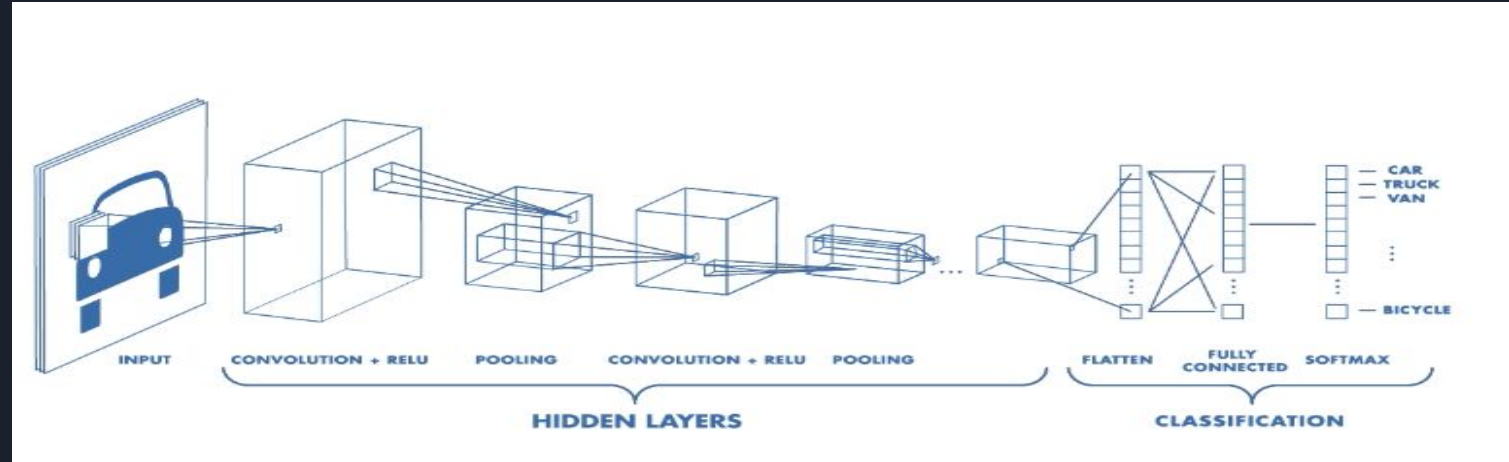
Type 1	Type 2	Type 3
250	781	450
17%	53%	30%

# Convolutional Neural Networks



Left: A regular 3-layer Neural Network. Right: A ConvNet arranges its neurons in three dimensions (width, height, depth), as visualized in one of the layers. Every layer of a ConvNet transforms the 3D input volume to a 3D output volume of neuron activations. In this example, the red input layer holds the image, so its width and height would be the dimensions of the image, and the depth would be 3 (Red, Green, Blue channels).

# Convolutional Neural Networks



- **Convolutional Layer:**  $n \times n \times d$  filters which pass through the image, and compute various “activations”
- **ReLU:** Adds non-linearity to model, with reduced problem of ‘vanishing gradients’
- **Pooling:** Downsample layer outputs to reduce computational load and overfitting (common methods include Maxpool, average pool)
- **Flatten:** Converts Output of final convolutional layer to 1-D Tensor (for fully connected layer).
- **Global Average Pooling:** Reduces computational overload of Flatten by averaging the output of final convolution (e.g  $20 \times 20 \times 15 \rightarrow 1 \times 1 \times 15$ )
- **Fully-Connected:** Perceptron to every pixel in output of final convolution
- **Softmax:** Multi-class generalization of the the sigmoid non-linearity function
- **Output**

# Prior Work

Payette, Rachleff, Van de Graaf (2018).












*Intel and MobileODT Cervical Cancer*

*Screening Kaggle Competition: Cervix Type*

*Classification Using Deep Learning and Image Classification, 2018*

- 32 Layer Residual CNN
- 32 Layer Residual CNN with Dropout and Augmentation
- 32 Layer Residual CNN with Augmentation
- 32 Layer Residual CNN with Dropout
- Validation Accuracy: 0.56 - 0.58
- Test Cross-Entropy Loss: 0.87 - 0.923

## Kaggle Leader Board

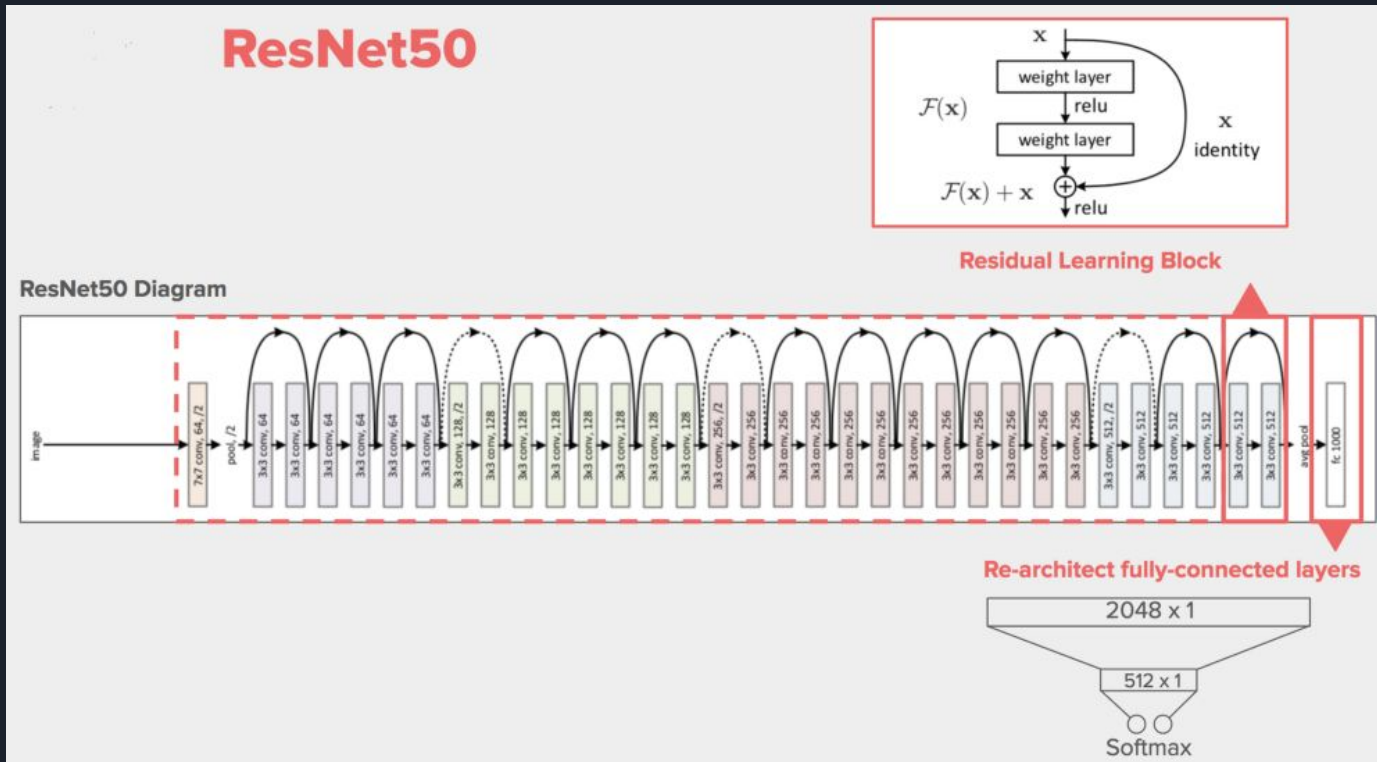
#	Δpub	Team Name	Kernel	Team Members	Score	Entries	Last
1	▲125	Towards Empirically Stable Tr...			0.76963	2	1y
2	▲114	I Rustandi			0.80277	8	1y
3	▲79	GRXJ			0.80829	5	1y
4	▲94	Ruslan Baikulov			0.81808	4	1y
5	▲73	BMC1			0.82205	5	1y
6	▲80	kubilai			0.82495	26	1y
7	▲124	Shai			0.82849	2	1y
8	▲93	Ryan Munion			0.83130	3	1y
9	▲78	ZFTurbo			0.83208	2	1y
10	▲134	Alexander Popov			0.83333	19	1y
11	▲136	mokp			0.83367	2	1y

Model	Val Loss	Val Accuracy	Val F1 Score	Test Loss
32 Layer Residual CNN	.948	.580	.573	0.923
32 Layer Residual CNN with Dropout	.948	.564	.567	-
32 Layer Residual CNN with Dropout and Data Augmentation	.936	.576	.551	-
32 Layer Residual CNN with Data Augmentation	.879	.588	.578	.873

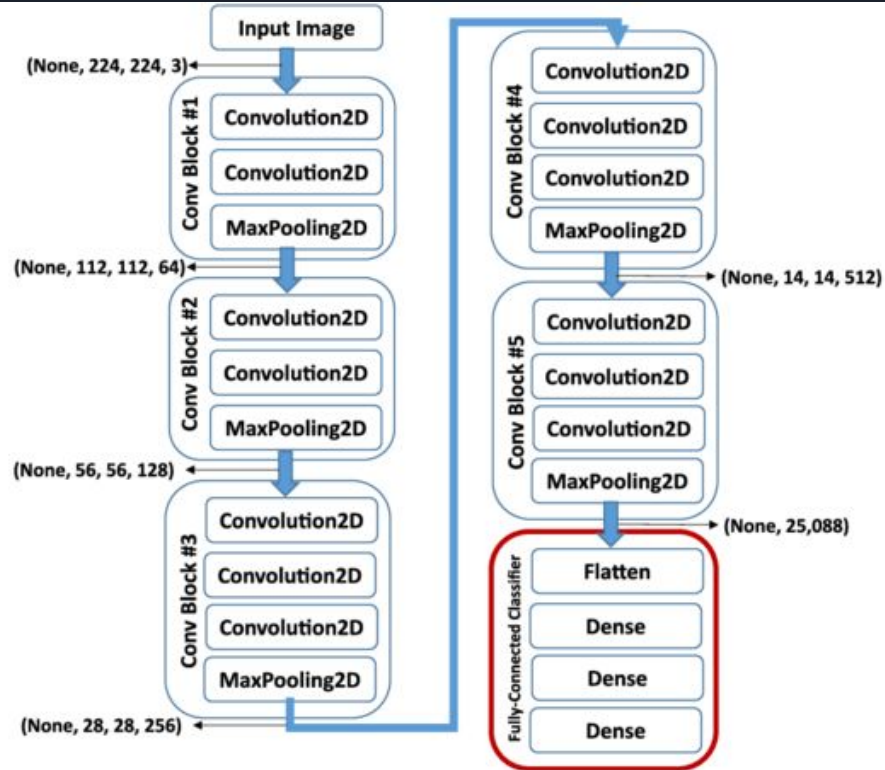
Steps followed to solve  
problem



# ResNet50 diagram



# VGG16 Inspired Model





# VGG16 Inspired Model. Keras implementation

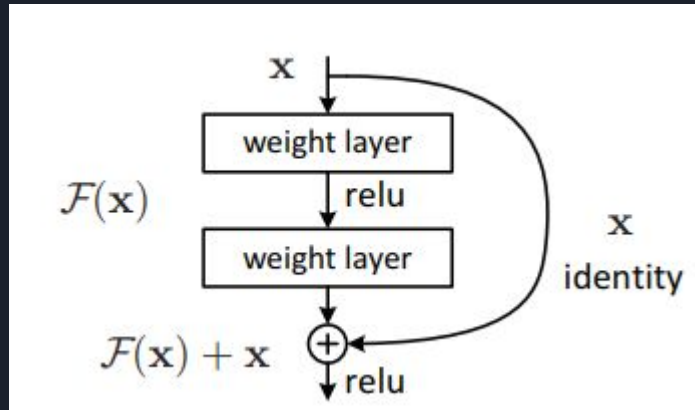
```
nClasses = 3
model = Sequential()
model.add(ZeroPadding2D((1, 1), input_shape=(64, 64, 3)))
model.add(Conv2D(64, kernel_size=3, strides = 1, activation='relu',
input_shape=input_shape))
model.add(ZeroPadding2D((1, 1)))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(ZeroPadding2D((1, 1)))
model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(ZeroPadding2D((1, 1)))
model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(ZeroPadding2D((1, 1)))
model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(ZeroPadding2D((1, 1)))
model.add(Conv2D(256, (3, 3), activation='relu'))
model.add(ZeroPadding2D((1, 1)))
model.add(Conv2D(256, (3, 3), activation='relu'))
model.add(Conv2D(256, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(ZeroPadding2D((1, 1)))
```

```
model.add(Conv2D(512, (3, 3), activation='relu'))
model.add(ZeroPadding2D((1, 1)))
model.add(Conv2D(512, (3, 3), activation='relu'))
model.add(ZeroPadding2D((1, 1)))
model.add(Conv2D(512, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(ZeroPadding2D((1, 1)))
model.add(Conv2D(512, (3, 3), activation='relu'))
model.add(ZeroPadding2D((1, 1)))
model.add(Conv2D(512, (3, 3), activation='relu'))
model.add(ZeroPadding2D((1, 1)))
model.add(Conv2D(512, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())
model.add(Dropout(0.2))
model.add(Flatten())
model.add(Dense(4096, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(4096, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(nClasses, activation='softmax'))
model.compile(loss='categorical_crossentropy',
optimizer='adam', metrics=['accuracy'])
```

# Model selection and implementation

## ResNet 50 Model

- Uses batch normalization
- Learns residuals and uses them to get an optimized output



$$\hat{x}^{(k)} = \frac{x^{(k)} - E[x^{(k)}]}{\sqrt{\text{Var}[x^{(k)}]}}$$

$$y^{(k)} = \gamma^{(k)} \hat{x}^{(k)} + \beta^{(k)}.$$

**Input:** Values of  $x$  over a mini-batch:  $B = \{x_1 \dots x_m\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_B \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_B^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} \quad // \text{ normalize}$$

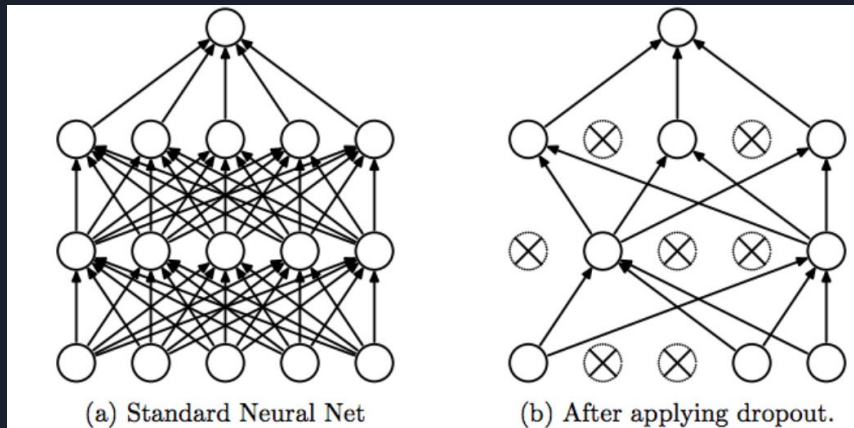
$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

**Algorithm 1:** Batch Normalizing Transform, applied to activation  $x$  over a mini-batch.

# Dropout Regularization

Keras Implementation:

```
keras.layers.Dropout(rate, noise_shape=None, seed=None)
```





# Training and testing

## Metrics

- Accuracy: Percentage of Correct Predictions
- Multi-class Cross-Entropy Loss:

$$\text{loss} = -\frac{1}{N} \sum_{i=1}^N \sum_{j=1}^M y_{ij} \ln(p_{ij})$$

## Train-Validation Split

- 75% - 25%
- Plain, randomized split

# Optimization

Improve training time

- Adam optimizer - Kingma & Ba (2015)
  - Improvement of the SGD optimizer and momentum optimizer
  - Computes first and second moments
  - Proven to converge faster than other optimizers and very widely used

**Require:**  $\alpha$ : Stepsize

**Require:**  $\beta_1, \beta_2 \in [0, 1]$ : Exponential decay rates for the moment estimates

**Require:**  $f(\theta)$ : Stochastic objective function with parameters  $\theta$

**Require:**  $\theta_0$ : Initial parameter vector

$m_0 \leftarrow 0$  (Initialize 1<sup>st</sup> moment vector)

$v_0 \leftarrow 0$  (Initialize 2<sup>nd</sup> moment vector)

$t \leftarrow 0$  (Initialize timestep)

**while**  $\theta_t$  not converged **do**

$t \leftarrow t + 1$

$g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$  (Get gradients w.r.t. stochastic objective at timestep  $t$ )

$m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$  (Update biased first moment estimate)

$v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$  (Update biased second raw moment estimate)

$\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$  (Compute bias-corrected first moment estimate)

$\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$  (Compute bias-corrected second raw moment estimate)

$\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$  (Update parameters)

**end while**

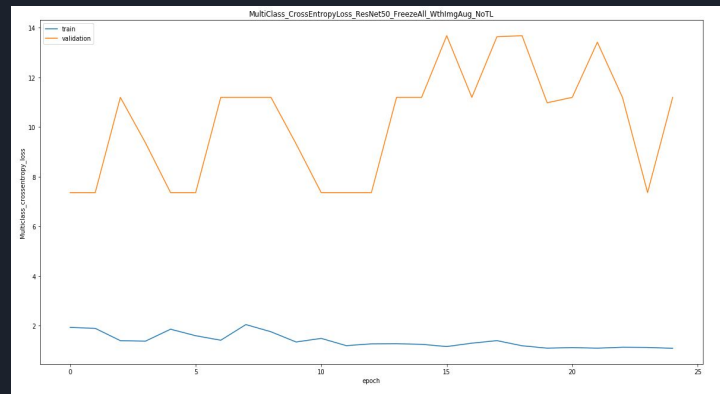
**return**  $\theta_t$  (Resulting parameters)

# Results

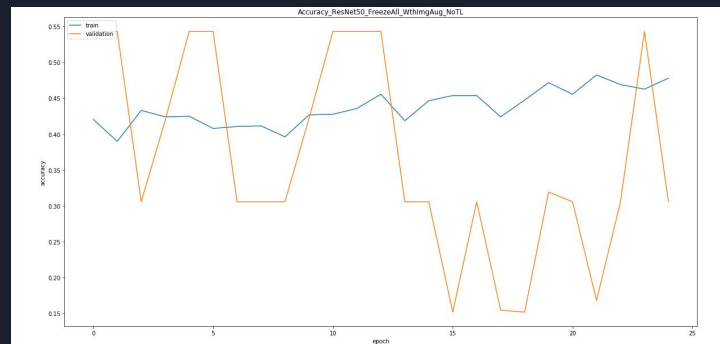
## No Transfer Learning

- Unpredictable and unstable behavior on validation data
- Validation multi-class Loss hovering around 8 - 12
- Validation accuracy 0.15 - 0.55

Multi-Class Entropy Loss

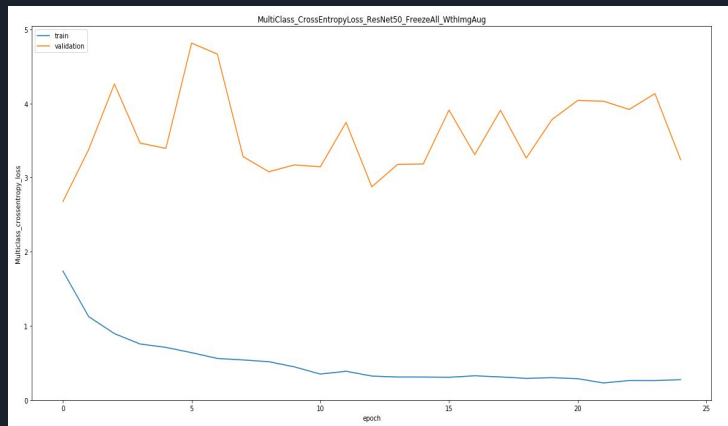


Accuracy

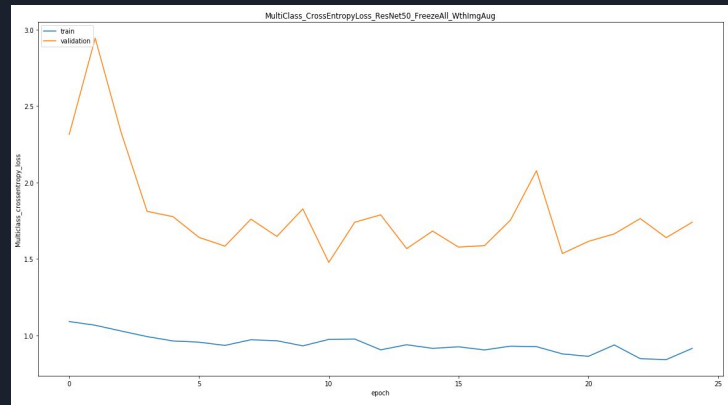


# Results - ResNet50

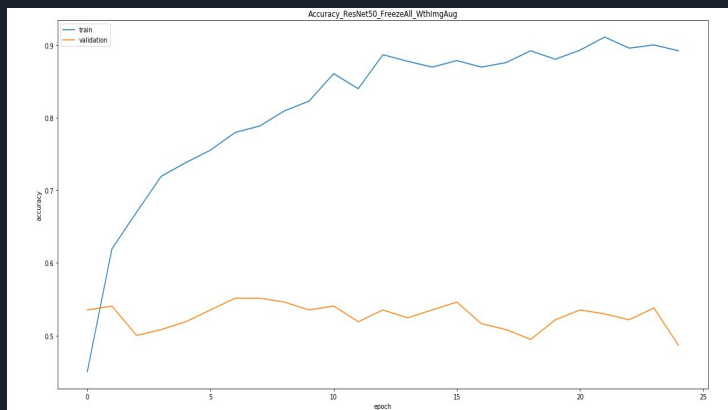
Multi-Class Entropy Loss No Augmentation



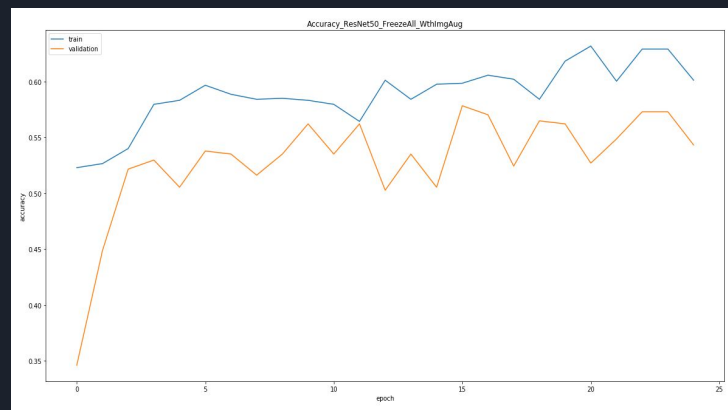
Multi-Class Entropy Loss With Augmentation



Accuracy No Augmentation

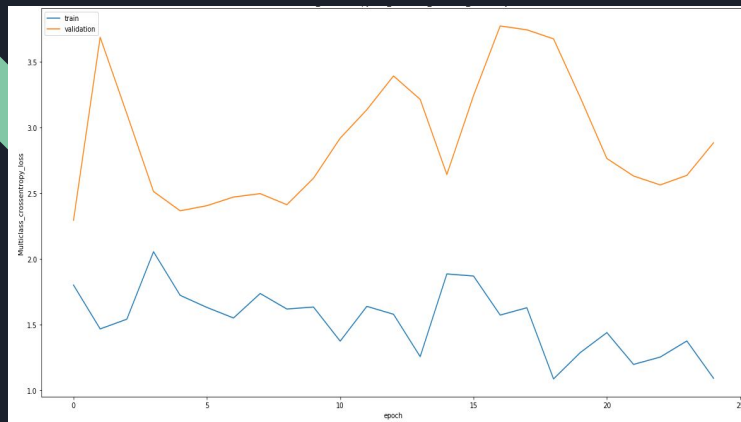


Accuracy With Augmentation



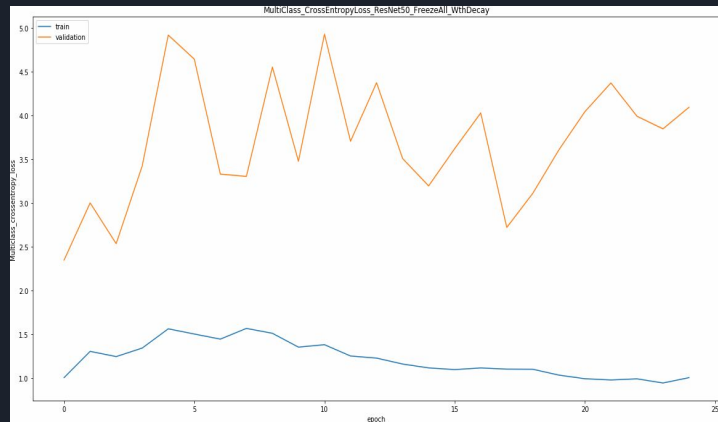
# Results - ResNet50

Multi-Class Entropy Loss Fixed Learning rate



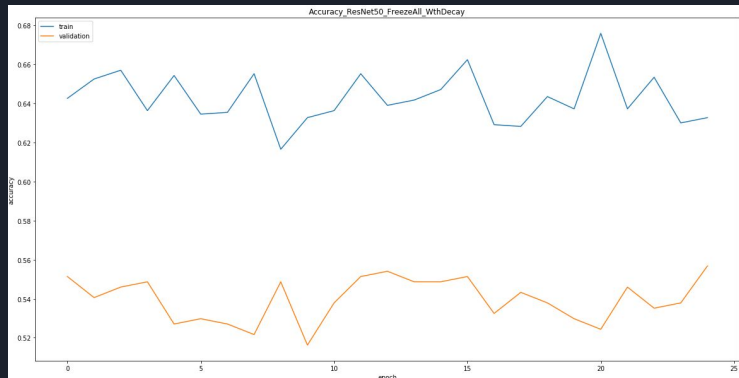
Loss:  
T: ~1.5  
V: 2.2 -  
3.6

Multi-Class Entropy Loss Decaying Learning rate



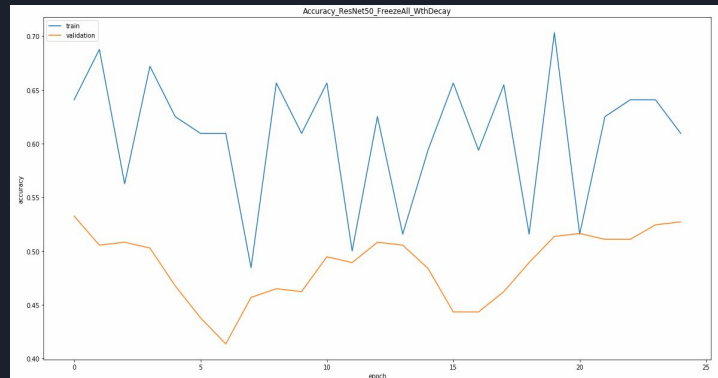
Loss:  
T: ~1.2  
V: 4

Accuracy Fixed Learning rate



Acc:  
T: ~0.64  
V: ~0.54

Accuracy Decaying Learning rate

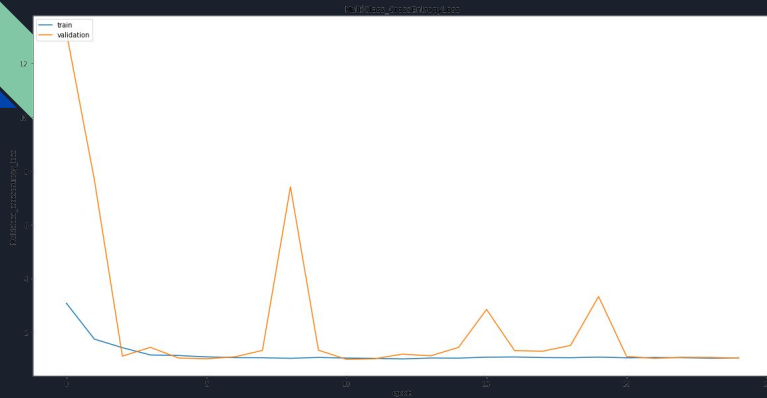


Acc:  
T: 0.55 -  
0.7  
V: 0.45 -  
0.55



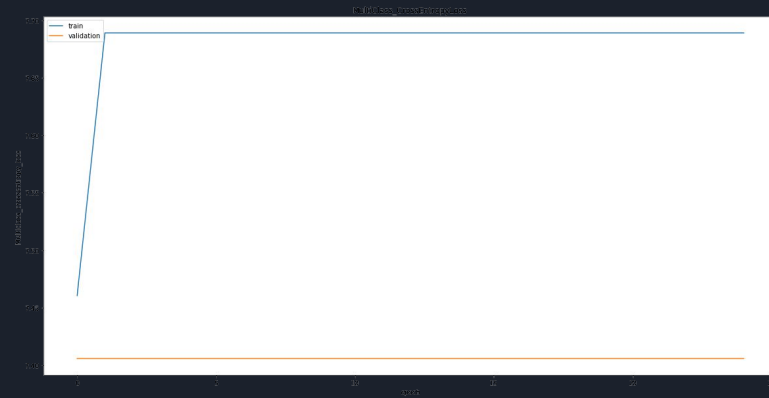
# Results - VGG16 Inspired

Multi-Class Entropy Loss w/ Batch Normalization & Dropout



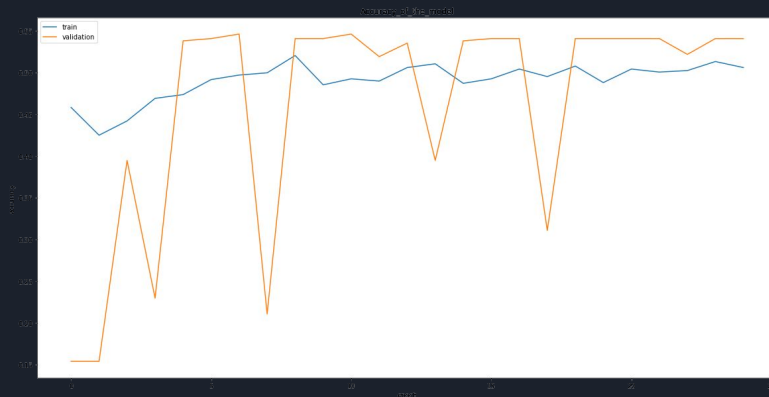
Loss:  
T: <1.0  
V: <1.0

Multi-Class Entropy Loss without Batch Normalization & Dropout



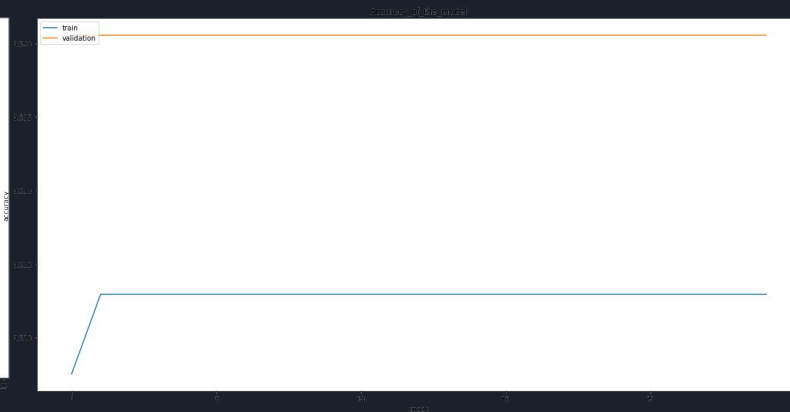
Loss:  
T: ~7.50  
V: ~7.41

Accuracy w/ Batch Normalization & Dropout



Acc:  
T: ~0.50  
V: ~0.55

Accuracy without Batch Normalization & Dropout

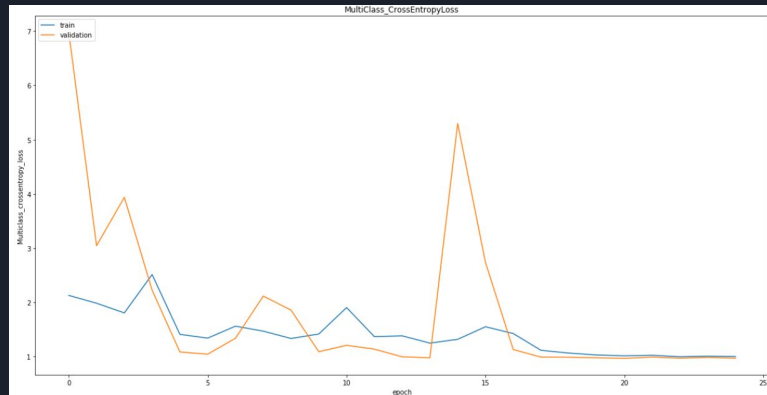


Acc:  
T: ~0.52  
V: ~0.54

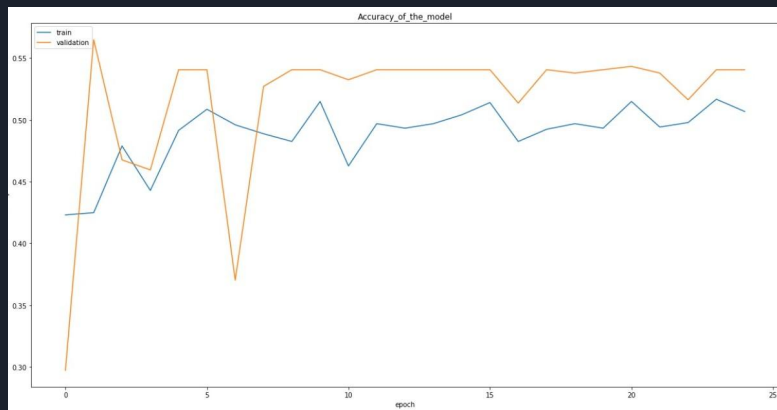
# Results - VGG16 Inspired

Multi-Class Entropy Loss w/ Batch Normalization & no Dropout

Loss:  
T: <1.0  
V: <1.0



Accuracy w/ Batch Normalization & no Dropout



Acc:  
T: ~0.50  
V: ~0.54

When you spend hours training your neural network and it finally gives you an accuracy of 37%



Am I a joke to you?

# Possible future enhancements

- Training with More Data (additional 4633 images were provided)
- Higher-Image Resolution: 224 x 224 vs 64 x 64
- Exploration of other Models (InceptionV3, Custom-built models)
- Ensembling of Models (Bagging, Boosting, Stacking)
- Enforcement of a more balanced class, confusion matrix could help with this

Predicted	Truth						
	Asphalt	Concrete	Grass	Tree	Building	Total	
	Asphalt	2385	4	0	1	4	2394
	Concrete	0	332	0	0	1	333
	Grass	0	1	908	8	0	917
	Tree	0	0	0	1084	9	1093
	Building	12	0	0	6	2053	2071
	Total	2397	337	908	1099	2067	6808

- Data augmentation: Payette, Rachleff, Van de Graaf (2018) utilize a method of statistical cropping. Sharpening. Embossing
- Is the problem entirely solvable by a machine?

Thank you! Questions?





# Sources

1. <https://www.kaggle.com/c/intel-mobileodt-cervical-cancer-screening>
2. <https://medium.freecodecamp.org/an-intuitive-guide-to-convolutional-neural-networks-260c2de0a050>
3. Payette, Rachleff, Van de Graaf (2018). *Intel and MobileODT Cervical Cancer Screening Kaggle Competition: Cervix Type Classification Using Deep Learning and Image Classification*. <http://cs231n.stanford.edu/reports/2017/pdfs/300.pdf>
4. Greenspan, Hayit & Gordon, Shiri & Zimmerman, Gali & Lotenberg, Shelly & Jeronimo, Jose & Antani, Sameer & Long, L. (2009). Automatic Detection of Anatomical Landmarks in Uterine Cervix Images. IEEE transactions on medical imaging. 28. 454-68. 10.1109/TMI.2008.2007823.
5. Kingma & Ba (2015). Adam: A Method for Stochastic Optimization. <https://arxiv.org/pdf/1412.6980.pdf>
6. Ioffe & Szegedy (2015). Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift <https://arxiv.org/pdf/1502.03167.pdf>
7. He, Zhang, Ren & Sun (2015). Deep Residual Learning for Image Recognition <https://arxiv.org/pdf/1512.03385.pdf>
8. [http://slazebni.cs.illinois.edu/spring17/lec01\\_cnn\\_architectures.pdf](http://slazebni.cs.illinois.edu/spring17/lec01_cnn_architectures.pdf)
9. <https://towardsdatascience.com/neural-network-architectures-156e5bad51ba>
10. [https://www.harrisgeospatial.com/docs/ENVIConfusionMatrix\\_ConfusionMatrix.html](https://www.harrisgeospatial.com/docs/ENVIConfusionMatrix_ConfusionMatrix.html)